

Computer Vision I

Assignment 5

Prof. Stefan Roth
Dustin Carrión
Jan-Martin Steitz

29/01/2024



TECHNISCHE
UNIVERSITÄT
DARMSTADT

This assignment is due on February 11th, 2024 at 23:59.

Please refer to the previous assignments for general instructions and follow the handin process described there.

Problem 1: Neural Networks (20 Points)

In this problem, we will design a multilayer perceptron (MLP) to classify bike and plane images. Image and label data are already loaded in training and validation splits. The neural network is represented by the class `Network`, which builds the network according to a list containing input, hidden and output dimensions.

- First, initialize the weights and biases of the neural network with `init_wb`. Biases should be initialized to zero and weights should be initialized randomly according to a normal distribution. The variance of the normal distribution from which the weights are drawn should be equal to $1/n_{i-1}$, where n_{i-1} is the input dimension of the i -th layer. (2 points)

The network will use ReLU activations and predict class probabilities for the binary problem using the sigmoid function.

- Implement the `relu` function and back propagate error derivatives in `relu_backward`. (1 + 2 points)
- Implement the `sigmoid` function and back propagate error derivatives in `sigmoid_backward`. (1 + 2 points)

Next, implement forward and backward propagation for the neural network:

- Implement the function `layer_forward` to perform forward propagation through a single layer. This function may make use of `activation_func`, which selects and calls an activation function. `layer_forward` should update the state dictionaries for `z[i]`, the output of the i -th linear layer, and `x[i]`, the output of the i -th activation function, since these values are required to compute the backward propagation. (3 points)
- Implement the function `forward`, which iteratively calls `layer_forward` in order to perform forward propagation and outputs the network's prediction. (2 points)
- Next, implement function `layer_backward` that propagates the derivatives backward through a layer. Here, the values stored in the state dictionaries `z[i]` and `x[i-1]` have to be used. This function may make use of `activation_func_backward`, which selects and calls the backward method of an activation function. The function `layer_backward` computes the partial derivatives of the error function with respect to the layer parameters `w[i]` and `b[i]` and updates the states of `dw[i]` and `db[i]` accordingly. Further, it returns the partial derivatives of the error function with respect to the layer's input. (3 points)
- We provide the function `back_propagation`, which iteratively calls `layer_backward` to propagate derivatives through the network in order to update the states of `dw` and `db` for all layers.

We further require the following methods to train the network:

- Implement the function `update_wb` to update the states of the network's weights `w` and biases `b` by performing a SGD step given the learning rate and states of `dw` and `db`. (2 points)
- Finally, implement the function `shuffle_data`, which permutes the training data so that we can extract random mini batches required for SGD from it. (2 points)
- The function `train` extracts mini batches from the shuffled arrays to train the neural network.

We already provide methods for evaluating the prediction accuracy of the neural network. You can try changing network and training parameters in `main.py` to improve the network's predictive performance.

Submission: Please include only `problem1.py` in your submission.

Problem 2: Convolutional Neural Networks (10 Points)

In this problem, we will solve the same task as in Problem 1 but using a convolutional neural network (CNN). We will implement a simple CNN architecture (inspired by VGG-11) using the PyTorch library (<https://pytorch.org/docs/stable/index.html>).

We will have four convolutional blocks for the feature extraction (blue modules in Figure 1). Each convolutional block contains a stack of layers; convolutional (Conv2D), batch normalization (BatchNorm), ReLU activation (ReLU) and max pooling (MaxPool2D). In order to use these features to solve the same classification task in Problem 1, we're using a two-layer Fully Connected (FC) classification head (orange modules in Figure 1). Finally, train the whole architecture using the cross-entropy loss.

For the implementation you may find the following functions from `torch.nn` helpful: `Conv2D()`, `BatchNorm2d()`, `ReLU()`, `MaxPool2d()`, `Linear()`, `Softmax()`.

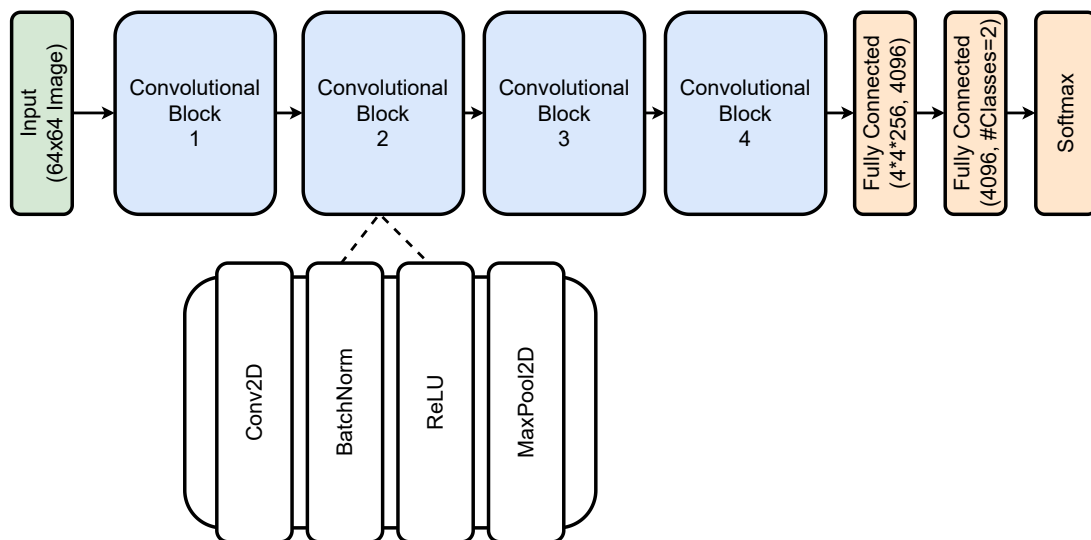


Figure 1: Convolutional Neural Network Architecture.

- First, implement the convolutional part of the model. Details required for the implementation are as follows:
(5 points)
 - Input-output dimensions of 2D convolutional blocks should be (1×64) , (64×128) , (128×256) , (256×256) , respectively.
 - You can set *kernel size*, *stride*, and *padding* parameters to 3, 1, 1, respectively, in each Conv2D layer.
 - You should set the input dimension of each BatchNorm2D layers by considering the output dimensions of the preceding Conv2Ds.
 - You can set the *kernel size* and *stride* parameters of each MaxPool2D layer to 2, 2, respectively.
- Implement the FC module of the network.
(3 points)
- Train the network with the given hyper-parameters in `main.py`. Compare the performance of the two networks in Problem 1 and Problem 2 and briefly summarize which one performs better and why. Include your comment in the function `analysis`.
(2 points)

Submission: Please include only `problem2.py` in your submission.